

# EXERCISE CLASS 3 (Part 2/3)

## EXERCISE 2

A study was performed by ComputerTek Co to determine the time series of order processing durations. Data in the file "dataset\_ese3\_es2.csv" refer to the period 1995, July – 1997, October. Each datum represents the time (in days) to ship the order.

Design a 95% **prediction interval** for future observations.

### Suggestion

Remind that if:

$$X \sim N(\mu, \sigma^2)$$

then:

$$\frac{X - \mu}{s} \sim t_{n-1}$$

### Confidence intervals

$L \leq \theta \leq U$  such that  $P(L \leq \theta \leq U) = 1 - \alpha$

is called  $100(1 - \alpha)\%$  confidence interval for the **(unknown) parameter  $\theta$**

*Interpretation:* if, in repeated random samplings, a large number of such intervals is constructed,  $100(1 - \alpha)\%$  of them will contain the true value of  $\theta$ .

### Prediction intervals

$L \leq Y \leq U$  such that  $P(L \leq Y \leq U) = 1 - \alpha$

is called  $100(1 - \alpha)\%$  prediction interval for the **future process outcome  $Y$**

*Interpretation:* if, in repeated random samplings, a large number of such intervals is constructed,  $100(1 - \alpha)\%$  of them will contain the future outcome  $Y$

```
In [ ]: # Import the necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from scipy import stats

# Import the dataset
data = pd.read_csv('dataset_ese3_es2.csv')
```

```
data.head()
```

Out[ ]: **Ex2**

0	29
1	21
2	17
3	9
4	49

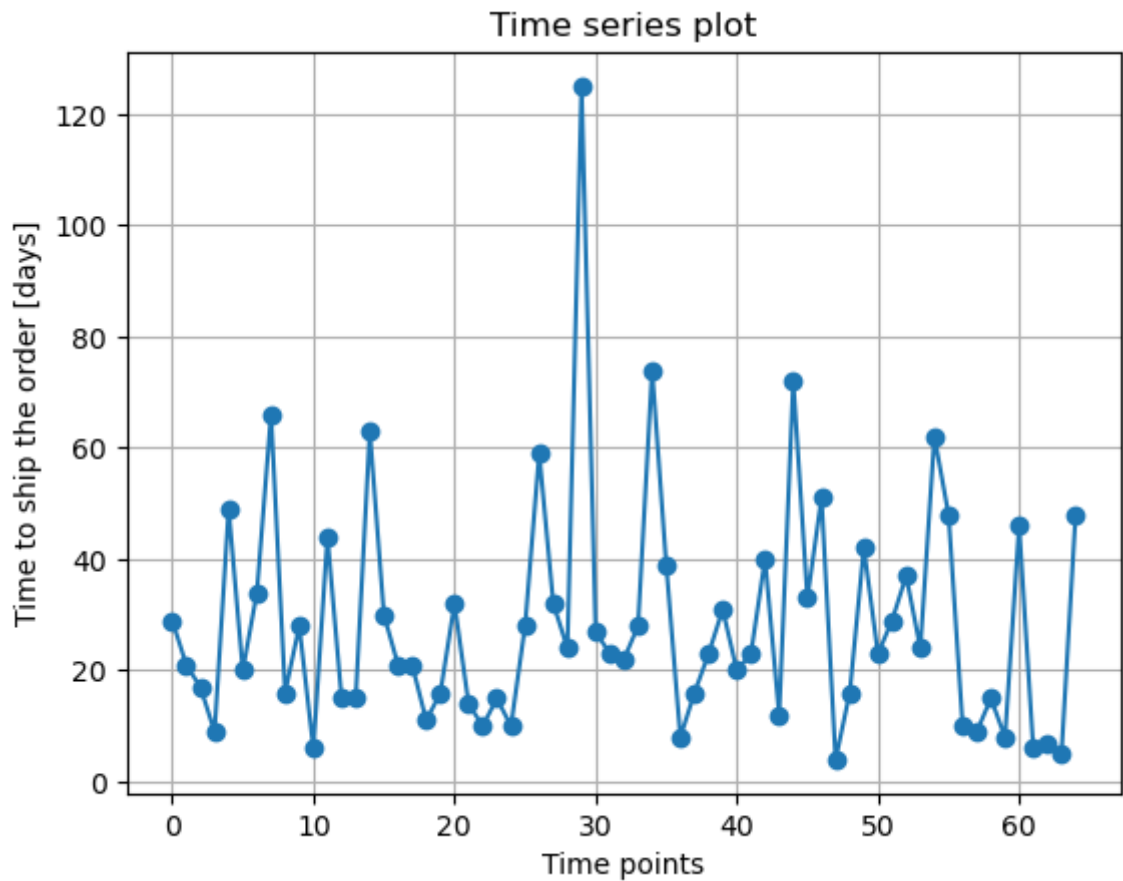
We need a model of process output to make predictions. Which kind of model? **Time series** model or **distributional** model.

We can verify if the process is random by using :

1. time series plot (qualitative)
2. ACF/PACF (qualitative)
3. runs test (quantitative)
4. Bartlett's test (quantitative)
5. LBQ test (quantitative)

1. Time series plot

```
In [ ]: # Time series plot
plt.plot(data, 'o-')
plt.xlabel('Time points')
plt.ylabel('Time to ship the order [days]')
plt.title('Time series plot')
plt.grid()
plt.show()
```



#### 1. Runs test

```
In [ ]: # Import the necessary libraries for the runs test
from statsmodels.sandbox.stats.runs import runtest_1samp


stat_runs, pval_runs = runtest_1samp(data['Ex2'], correction=False)
print('Runs test statistic = {:.3f}'.format(stat_runs))
print('Runs test p-value = {:.3f}'.format(pval_runs))
```

Runs test statistic = 0.325  
Runs test p-value = 0.745

#### 1. ACF/PACF

```
In [ ]: # Plot the acf using the statsmodels library
import statsmodels.graphics.tsaplots as sgt



sgt.plot_acf(data['Ex2'], lags = int(len(data)/3), zero=False)
plt.show()
```



#### 1. Histogram and Boxplot

```
In [ ]: plt.hist(data)
plt.title('Histogram')
plt.show()

plt.boxplot(data)
plt.title('Boxplot')
plt.show()
```



Let's check normality

```
In [ ]: #Normality test
#Shapiro-Wilk test
import matplotlib.pyplot as plt

stat_shapiro, p_shapiro = stats.shapiro(data)
```


```

print('Statistic = %.3f, p-val = %.3f' % (stat_shapiro, p_shapiro))
# interpret
alpha = 0.05
if p_shapiro > alpha:
    print('Fail to reject H0')
else:
    print('Reject H0')

# Plot the qqplot
stats.probplot(data['Ex2'], dist="norm", plot=plt)
plt.show()

```

Statistic = 0.844, p-val = 0.000  
Reject H0



How much is this result influenced by the outlier? We can try to remove the outlier and check for normality again.

```

In [ ]: # Remove outlier (point 30) and check normality
data_out = data.drop(index=29)

plt.hist(data_out)
plt.title('Histogram')
plt.show()


plt.boxplot(data_out)
plt.title('Boxplot')
plt.show()


#Normality test
#Shapiro-Wilk test
from scipy.stats import shapiro
stat_shapiro_out, p_shapiro_out = shapiro(data_out)

```


```
print('Statistic = %.3f, p-val = %.3f' % (stat_shapiro_out, p_shapiro_out))
# interpret
alpha = 0.05
if p_shapiro > alpha:
    print('Fail to reject H0')
else:
    print('Reject H0')

# Plot the qqplot
stats.probplot(data_out['Ex2'], dist="norm", plot=plt)
plt.show()
```





Statistic = 0.914, p-val = 0.000  
Reject  $H_0$



Even after removing the outlier, normality is still violated.

Try with the Box-Cox transformation.



## Remind:

$$y = (x^\lambda - 1)/\lambda, \text{ for } \lambda \neq 0$$


$$y = \ln(x), \text{ for } \lambda = 0$$

```
In [ ]: #Box-Cox transformation
[data_norm, lmbda]=stats.boxcox(data['Ex2'])

print('Lambda = %.3f' % lmbda)

plt.hist(data_norm)
plt.title('Histogram of Box-Cox transformed data')
plt.show()
```


Lambda = 0.100



```
In [ ]: stat, p_shapiro = stats.shapiro(data_norm)
print('Statistics=%.3f, p=%.3f' % (stat, p_shapiro))
# interpret
alpha = 0.05
if p_shapiro > alpha:
    print('Fail to reject H0')
else:
    print('Reject H0')


# Plot the qqplot
stats.probplot(data_norm, dist="norm", plot=plt)
plt.show()
```

Statistics=0.992, p=0.959  
Fail to reject H0



```
In [ ]: # Plot the acf using the statsmodels library

sgt.plot_acf(data_norm, lags = int(len(data_norm)/3), zero=False)
plt.show()
```



No graphical evidence of auto-correlation from the sample acf plot. We can verify it with quantitative tests, e.g.:

- Bartlett's test
- LBQ test

On transformed data

### Bartlett's test for a specific lag $k$

$$H_0 : \rho_k = 0$$

$$H_1 : \rho_k \neq 0$$

$\rho_k$ : true autocorr at lag  $k$

$r_k$ : sample autocorr at lag  $k$

$$\text{Rejection region } |r_k| > \frac{z_{\alpha/2}}{\sqrt{n}}$$

```
In [ ]: from statsmodels.tsa.stattools import acf

n = len(data_norm)

#autocorrelation function
[acf_values, lbq, _] = acf(data_norm, nlags = int(np.sqrt(n)), qstat=True, fft = F
```

```
In [ ]: #Bartlett's test at lag 1
lag_test = 1
rk = acf_values[lag_test]
print('Test statistic rk = %f' % rk)
print('Rejection region starts at %f' % (1.96/np.sqrt(n)))

if rk>1.96/np.sqrt(n):
    print('The null hypothesis is rejected')
else: print('The null hypothesis is accepted')
```

Test statistic rk = 0.040639  
 Rejection region starts at 0.243108  
 The null hypothesis is accepted

### LBQ test

$$H_0 : \rho_k = 0, k = 1, \dots, L$$

$$H_1 : \exists k \in [1, L] \text{ such that } \rho_k \neq 0$$

LBQ test statistic:

$$LBQ = n(n+2) \sum_{k=1}^L \frac{r_k^2}{n-k}$$

Under  $H_0$  ( $\rho_k = 0, k = 1, \dots, L$ ),  $LBQ \sim \chi_{L'}^2$ , and its rejection region is:

$$LBQ > \chi_{\alpha, L}^2$$

```
In [ ]: lag_test=6 # this is just an example;
```

```
# Generally speaking: how many lags?
# Rule of thumb:  $L < \sqrt{n}$ 

Q0_LBQ = lbq[lag_test-1]
print('Q0_LBQ = %f' % Q0_LBQ)

#Rejection region for chi square distribution
dof = lag_test
chi2_alfa= stats.chi2.ppf(1-alpha,dof)
print('Rejection region starts at %f' % chi2_alfa)

if Q0_LBQ>chi2_alfa:
    print('The null hypothesis is rejected')
else:
    print('The null hypothesis is accepted')

# Compute the p-value for the LBQ test
pval = 1 - stats.chi2.cdf(Q0_LBQ, lag_test)
print('p-value = %f' % pval)

Q0_LBQ = 0.621926
Rejection region starts at 12.591587
The null hypothesis is accepted
p-value = 0.996024
```

Alternatively, you can use the `acorr_ljungbox` function.

```
In [ ]: #LBQ test for autocorrelation
from statsmodels.stats.diagnostic import acorr_ljungbox

lbq_test = acorr_ljungbox(data_norm, lags=[lag_test], return_df=True)
print('LBQ test statistic at lag %d = %f' % (lag_test, lbq_test.loc[lag_test,'lb_stat'])
print('LBQ test p-value at lag %d = %f' % (lag_test, lbq_test.loc[lag_test,'lb_pval'])

LBQ test statistic at lag 6 = 0.621926
LBQ test p-value at lag 6 = 0.996024
```

...to finally answer the question (Design a 95% prediction interval for future observations):

Process data are **normal and independent (NID)**

$$X \sim N(\mu, \sigma^2) \rightarrow \frac{X - \mu}{s} \sim t_{n-1} \quad \text{Attention: this is not } T = \frac{\bar{X} - \mu}{S/\sqrt{n}}$$

Prediction interval:

$$\Pr\left(-t_{\alpha/2, n-1} \leq \frac{X - \mu}{s} \leq t_{\alpha/2, n-1}\right) = 1 - \alpha \quad \Rightarrow \quad \mu - t_{\alpha/2, n-1}s \leq X \leq \mu + t_{\alpha/2, n-1}s$$

**Approximated** prediction interval (95%): Approximated because we use the sample mean in place of the true mean

$$\bar{X} \pm t_{0.025, 64} \cdot s$$

```
In [ ]: alpha = 0.05
df = len(data_norm) - 1
Xbar = data_norm.mean()
s = data_norm.std()
```

```
t_alpha = stats.t.ppf(1 - alpha/2, df)
```

```
[pred_lo, pred_up] = [Xbar-t_alpha*s, Xbar+t_alpha*s]
```

```
print('Two-sided confidence interval for transformed data: [%.3f %.3f]' % (pred_lo,
```

Two-sided confidence interval for transformed data: [1.719 5.640]

**Attention:** this is the prediction interval on the transformed data. To estimate the prediction interval on the **original data** we need to back transform.

Remind:  $y = (x^\lambda - 1)/\lambda$ , for  $\lambda \neq 0$

Thus:  $x = (y\lambda + 1)^{(1/\lambda)}$

```
In [ ]: [pred_lo_ORIG, pred_up_ORIG] = [(pred_lo*lmbda+1)**(1/lmbda), (pred_up*lmbda+1)**(1,
print('Two-sided confidence interval for original data: [%.3f %.3f]' % (pred_lo_ORIG,
```

Two-sided confidence interval for original data: [4.887 87.689]